

Pujan Paudel

 pujanpaudel.com  ppaudel@bu.edu  Boston, MA  725-724-0359  [codepujan](#)

EDUCATION

Boston University (BU)

Boston, MA

Ph.D. in Computer Engineering

September 2025 (expected)

M.S. in Computer Engineering, GPA: 3.93

December 2023

Research Areas: Natural Language Processing (NLP), Machine Learning (ML), Deep Learning (DL), Information Retrieval (IR), Computer Vision (CV), and Cybersecurity

The University of Southern Mississippi (USM)

Hattiesburg, MS

B.S. (Honors) in Computer Science, GPA: 3.75

May 2020

Honors Thesis: *The Blind Spot of Twitter Bot Moderation*

WORK EXPERIENCE

Machine Learning (ML) PhD Intern at **Truveta**

September 2024 - December 2024

- Developed an end-to-end Retrieval Augmented Generation (RAG) system for automatic semantic normalization of unstructured clinical data to standard medical ontologies.
- Improved term mapping coverage by 36% while deploying RAG system across Truveta's large-scale Electronic Health Record (EHR) datasets.

Natural Language Processing (NLP) Research Intern at **Merck**

June 2024 - August 2024

- Curated domain-specific biomedical text corpus, benchmarked state-of-the-art NLP models and developed a hybrid BERT-LLM pipeline for automated document classification.
- Integrated the document classifier pipeline on proprietary stakeholder data, improving identification of target biomedical documents by up to 8%.

Graduate Research Assistant at **Boston University Security Lab**

September 2020 - Present

- Proactively discovering e-commerce scam websites for upstream online fraud mitigation [1].
 - ❑ Conceptualized and developed *Loki*, a feature distillation-based framework for Search Engine Ranking Pages (SERP) driven query mining system.
 - ❑ Expanded discovery of scam websites by 21 times relative to baseline NLP models, accelerating discovery of emergent online scams in the wild.
- Contextual content moderation system for precise flagging of harmful content [3].
 - ❑ Developed and implemented a new task for unsupervised stance detection leveraging transfer learning and fine-tuning Google's FLAN-T5 model.
 - ❑ Achieved state-of-the-art stance detection performance across 3 diverse social media datasets reducing false positives of automated content flagging 10 folds better.
- Identification of harmful media at Twitter for multimodal content moderation [4].
 - ❑ Implemented *PixelMod*, an end-to-end reverse image search system for million-scale social media image search using perceptual hashing and Milvus.
 - ❑ Improved precision and relevance of visual similarity matching by 8%, enhancing visual content flagging by 13 times on Twitter.
- Context adaptive keyword extraction system to track tweets spreading harmful content [5].
 - ❑ Developed *Lambretta*, a supervised Learning To Rank (LTR) and Open Information Extraction (OIE) based keyword identification system.
 - ❑ Improved the flagging of harmful content by 20 times on Twitter over a dataset of 4M tweets.

TECHNICAL SKILLS

Programming: Python, C, C++, JavaScript, Java, C#, Bash scripting, R

Machine Learning: Scikit-learn, PyTorch, TensorFlow, Pandas, Weights and Biases, OpenCV, SciPy, NumPy

Technologies: Elasticsearch, Lucene, AWS, Azure, GCP, Docker, Flask, FastAPI, Selenium, Spark, React.js, Django, Node.js, vLLM, Ollama, LangChain, Ray, Github Actions, Azure DevOps, Jenkins, MLflow

Database: Milvus, FAISS, Cassandra, MongoDB, SQL, PostgreSQL

Concepts: Data mining, Exploratory Data Analysis, Information Retrieval, Knowledge Distillation, Knowledge Graphs, Generative AI, Transformers, Computer Vision, Large Language Models, Multi-Modal LLMs, Recommender Systems, Search Systems, Human-Computer Interaction, Big Data, Cloud Computing

ADDITIONAL PROJECTS

American Rescue Plan Act (ARPA) Bill Earmark Analysis.

March 2022 - May 2022

- Built a custom Named Entity Recognition (NER) model using spaCy to automatically infer policy buckets from ARPA amendment language, analyzing disproportionate distribution of earmarked funding.

Scaling Remote Sensing Data Processing With Ray.

September 2021 - December 2021

- Setup OpenTelemetry, Ray and Jaeger in Mass Open Cloud (MOC) for distributed profiling and identifying bottlenecks on a NASA-JPL remote sensing application, optimizing a new parallelization scheme with 3x speedup.

Data-Driven Analysis of Targeted Cyberattacks and Effects on Foreign Policy.

May 2021 - July 2021

- Crawled, aggregated, and normalized a dataset of state-sponsored cyber attacks from three heterogeneous data sources, analyzing the change in severity of future cyberattacks as an effect of policy actions between rival nations.

Geolocation-based Food Recommendation System.

March 2021 - May 2021

- Benchmarked indexing and retrieval performance of KD-Tree and R-Tree based spatial indexes on 424K POIs, accompanied by an interactive companion web app that recommends nearby restaurants with multi-faceted search.

911 Overflow Management.

September 2019 - October 2019

- Developed an IVR system to relieve 911 call surges during disasters, integrating Twilio for scalable VOIP routing and Dialogflow for NLP-based intent extraction, coupled with a web-app to aggregate related incident reports.

Voice Agent based Tutor for Mathematics Retention.

August 2017 - February 2018

- Implemented a conversational agent using Amazon Alexa skills, assisting USM Psychology department in a study using conversational AI in retention of single-digit mathematical calculations for children with cognitive challenges.

AWARDS AND HONORS

Finalist, Best Applied Research Award, CSAW	2024
Usenix Security '24 Travel Grant, Usenix Association	2024
Pardee Center Graduate Summer Fellowship, BU	2021
Distinguished ECE PhD Fellowship, BU	2020
Runner Up, Undergraduate Research Symposium, USM	2019
Best Innovation Application, CalHacks	2016
Best IBM Watson Hack, HackRice	2016

SELECTED PUBLICATIONS

- [1] **P. Paudel**, G. Stringhini, "Loki: Proactively discovering online scams by mining toxic search queries," 35th Network and Distributed System Security Symposium (NDSS), 2026.
- [2] M.H. Saeed, S. Ali, **P. Paudel**, J. Blackburn, G. Stringhini, "Unraveling the Web of Disinformation: Exploring the Larger Context of State-Sponsored Influence Campaigns on Twitter," 27th International Symposium on Research in Attacks, Intrusions and Defenses (RAID 2024), Padua, Italy, 2024.
- [3] **P. Paudel**, M.H. Saeed, R. Auger, C. Wells and G. Stringhini, "Enabling Contextual Soft Moderation on Social Media through Contrastive Textual Deviation," 33rd Usenix Security Symposium, Philadelphia, PA, USA, 2024.
- [4] **P. Paudel**, C. Ling, J. Blackburn and G. Stringhini, "PixelMod: Improving Soft Moderation of Visual Misleading Information on Twitter," 33rd Usenix Security Symposium, Philadelphia PA, USA, 2024.
- [5] **P. Paudel**, J. Blackburn, E. De Cristofaro, S. Zannettou and G. Stringhini, "Lambretta: Learning To Rank For Twitter Soft Moderation," 2023 IEEE Symposium on Security and Privacy (SP), San Francisco, CA, USA, 2023.
- [6] M. Singhal, C. Ling, **P. Paudel**, P. Thota, N. Kumarswamy, G. Stringhini, and S.Nilizadeh "SoK: Content Moderation in Social Media, from Guidelines to Enforcement, and Research to Practice," 2023 IEEE 8th European Symposium on Security and Privacy (EuroS&P), 2023.
- [7] **P. Paudel**, TT. Nguyen, A. Hatua and AH. Sung, "How the tables have turned: Studying the new wave of social bots on Twitter using complex network analysis techniques," 2019 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining, Vancouver, Canada, 2019.
- [8] **P. Paudel**, TT. Nguyen, A. Hatua and AH. Sung, "User Level Multi-feed Weighted Topic Embeddings for Studying Network Interaction in Twitter." 2019: 8th International Congress, Held as Part of the Services Conference Federation, SCF 2019, San Diego, CA, USA, 2019.